

Practical-1

AIM: Study Of Machine learning basics

1. What is Machine learning

SOLUTION:

Machine learning is a part of artificial intelligence that allows computers to learn from data and improve their performance on a task without being explicitly programmed. It's like teaching machines to learn and make decisions based on experience.

2. Steps in collection of data

SOLUTION:

1. **Define Objectives:** Clearly outline the goals of data collection. What insights or information are you seeking?
2. **Plan and Design:** Determine the type of data needed and how it will be collected. Design surveys, experiments, or other methods accordingly.
3. **Select Data Sources:** Identify where to find the required data. This could involve surveys, existing databases, sensors, or other sources.
4. **Data Collection:** Execute the planned methods to gather data. This might involve surveys, observations, interviews, or automated sensors, depending on the nature of the study.
5. **Data Validation:** Ensure the collected data is accurate and reliable. This step involves checking for errors and inconsistencies.
6. **Data Entry:** If applicable, enter the data into a computer system for analysis. This step is crucial for manual data collection methods.
7. **Data Cleaning:** Identify and rectify any errors or inconsistencies in the collected data. This ensures the data is ready for analysis.
8. **Data Storage:** Organize and store the data securely. This step is critical for maintaining data integrity and privacy.

3. Steps in importing the data in python (Through: csv, json, and other data formats)

SOLUTION:

```
import csv # For CSV files
import json # For JSON files
import pandas as pd # For various data formats
```

Import CSV Data:

```
with open('your_file.csv', 'r') as file:
    csv_reader = csv.reader(file)
    data = [row for row in csv_reader]
```

Import JSON Data:

```
with open('your_file.json', 'r') as file:
    data = json.load(file)
```

For CSV, Excel, JSON, and more

```
data_csv = pd.read_csv('your_file.csv')
data_excel = pd.read_excel('your_file.xlsx')
data_json = pd.read_json('your_file.json')
```

```
# Display the first few rows of a pandas DataFrame
print(data_csv.head())
print(data_excel.head())
print(data_json.head())
```

4. Preprocessing

1. Remove Outliers

SOLUTION:

```
import pandas as pd
from scipy import stats
# Assuming 'data' is a pandas DataFrame
z_scores = stats.zscore(data)
filtered_data = data[(z_scores < 3).all(axis=1)] # Keep data points within 3 standard deviations
```

2. Normalize Datasets, Data encoding

SOLUTION:

```
from sklearn.preprocessing import MinMaxScaler
scaler = MinMaxScaler()
normalized_data = scaler.fit_transform(data)
```

DATA ENCODING:

```
# Assuming 'data' has categorical columns to be encoded
encoded_data = pd.get_dummies(data)
# Alternatively, you can use Label Encoding for ordinal categorical data
# from sklearn.preprocessing import LabelEncoder
# label_encoder = LabelEncoder()
# encoded_data['categorical_column'] = label_encoder.fit_transform(data['categorical_column'])
```

3. Handling Missing Data

SOLUTION:

```
# Drop rows with missing values
data_without_missing = data.dropna()
# Fill missing values with the mean of the column
data_filled_mean = data.fillna(data.mean())
```

5. Machine Models

1. Types of machine learning models – Supervised learning, Unsupervised learning, reinforcement learning.

SOLUTION:

a. Supervised Learning:

Definition: In supervised learning, the model is trained on a labeled dataset, where the input data is paired with corresponding target labels.

Example: Classification and Regression.

b. Unsupervised Learning:

Definition: Unsupervised learning involves training models on unlabeled data, and the algorithm tries to learn the patterns and structure within the data.

Example: Clustering and Dimensionality Reduction.

c. Reinforcement Learning:

Definition: Reinforcement learning involves an agent learning how to behave in an environment by performing actions and receiving rewards or penalties in return.

Example: Training a computer program to play a game.

2. Parameters of machine learning model (Learning rate, regularization, etc.)

SOLUTION:

Learning Rate:

The learning rate is crucial, influencing how quickly a model converges and its overall performance. If set too high, the model might overshoot the optimal solution; if too low, convergence may be slow.

Regularization:

Regularization combats overfitting by adding a penalty term to the loss function. L1 regularization (Lasso) and L2 regularization (Ridge) are common methods.

Hyperparameters:

These are pre-training external settings, not learned from data. Examples include the number of hidden layers in a neural network or the number of trees in a random forest.

Activation Function:

Activation functions like ReLU, Sigmoid, and Tanh introduce non-linearity, influencing the model's capacity to learn complex patterns.

Loss Function:

The loss function measures the disparity between predicted and actual values, guiding model training. For regression, Mean Squared Error (MSE); for classification, Cross-Entropy is often used.

6. Test-train data split: using constant ration, k-fold cross validation

SOLUTION:

a. Using Constant Ratio:

Divide the dataset into two parts, typically 70-30 or 80-20 for training and testing, respectively. The larger portion is used for training the model, and the smaller one is reserved for evaluating its performance. This simple approach provides a quick assessment but may lead to variability in results depending on the random split.

b. K-Fold Cross-Validation:

1. **Divide Data into K Folds:**

- Split the dataset into K equally-sized folds.

2. **Iterate Through Folds:**

- For each fold, designate it as the test set and the remaining K-1 folds as the training set.

3. **Train and Evaluate:**

- Train the model on the training set and evaluate its performance on the test set.

4. **Average Results:**

- Repeat this process K times, using a different fold as the test set in each iteration.
- Average the performance metrics to get a more reliable estimate of the model's performance.

7. Output Inference

SOLUTION:

1. **Prediction Interpretation:**

- Examine the model's predictions in the context of the problem. Understand the meaning and implications of the predicted values.

2. **Thresholding (if applicable):**

- For classification tasks, apply a threshold to convert raw model outputs into class predictions. This step is essential when dealing with probability scores.

3. **Evaluate Confidence:**

- Assess the model's confidence in its predictions. Some models provide uncertainty estimates or confidence intervals.

4. **Error Analysis:**

- Investigate prediction errors. Identify patterns or specific cases where the model struggles, helping improve model performance.

5. **Business/Application Context:**

- Consider the broader business or application context. How do the model's predictions impact decision-making or actions?

6. **Communication of Results:**

- Clearly communicate the model's outputs and any associated uncertainties to relevant stakeholders. Effective communication ensures informed decision-making.

7. **Iterative Improvement:**

- If applicable, use feedback and new data to iteratively improve the model. Continuous learning and refinement contribute to better performance over time.

8. **Decision Making:**

- Based on the model's predictions, make informed decisions or take appropriate actions in line with the problem's objectives.

8. Validation: different metrics – Confusion Matrix, Precision, Recall, F1-score

SOLUTION:

1. Confusion Matrix:

- A table that summarizes the performance of a classification model.
- **Components:**
 - **True Positive (TP):** Correctly predicted positive instances.
 - **True Negative (TN):** Correctly predicted negative instances.
 - **False Positive (FP):** Incorrectly predicted as positive.
 - **False Negative (FN):** Incorrectly predicted as negative.

2. Precision:

- Precision is the ratio of correctly predicted positive observations to the total predicted positives.
- **Formula:** $\text{Precision} = \text{TP} / (\text{TP} + \text{FP})$
- A high precision indicates a low false positive rate.

3. Recall (Sensitivity):

- Recall is the ratio of correctly predicted positive observations to the all observations in actual class.
- **Formula:** $\text{Recall} = \text{TP} / (\text{TP} + \text{FN})$
- A high recall indicates a low false negative rate.

4. F1-Score:

- F1-Score is the harmonic mean of precision and recall, providing a balanced evaluation metric.
- **Formula:** $\text{F1-Score} = 2 * (\text{Precision} * \text{Recall}) / (\text{Precision} + \text{Recall})$
- Particularly useful when there is an uneven class distribution.

Application:

- Use these metrics to comprehensively evaluate a classification model's performance, especially in scenarios where certain errors (false positives or false negatives) carry more significance.
- Consider the specific requirements of your problem to choose the most appropriate metric for model validation.